

# A2\_WriteUp

February 22, 2020

## 0 Setup

```
[1]: try:
      import google.colab
      IN_COLAB = True
    except:
      IN_COLAB = False
      import tensorflow as tf

    if IN_COLAB:
      %tensorflow_version 1.13
    else:
      assert tf.__version__ == "1.13.1"

    # ignore all future warnings
    from warnings import simplefilter
    simplefilter(action='ignore', category=FutureWarning)
```

`%tensorflow\_version` only switches the major version: `1.x` or `2.x`.  
You set: `1.13`. This will be interpreted as: `1.x`.

TensorFlow 1.x selected.

```
[0]: # imports
      import tensorflow as tf
      import numpy as np
      import matplotlib.pyplot as plt
      from tensorflow.keras import layers, models
```

```
[3]: print(tf.__version__)
```

1.15.0

```
[0]: # ignore tensorflow depreciation warnings
      import tensorflow.python.util.deprecation as deprecation
      deprecation._PRINT_DEPRECATION_WARNINGS = False
```

### 0.1 Visualizing the Dataset

```
[0]: # given by the assignment
      def loadData():
          with np.load("notMNIST.npz") as data:
              Data, Target = data["images"], data["labels"]
              np.random.seed(521)
              randIndx = np.arange(len(Data))
              np.random.shuffle(randIndx)
```


```
Data = Data[randIndx]/255.  
Target = Target[randIndx]  
trainData, trainTarget = Data[:15000], Target[:15000]  
validData, validTarget = Data[15000:16000], Target[15000:16000]  
testData, testTarget = Data[16000:], Target[16000:]  
return trainData, validData, testData, trainTarget, validTarget, testTarget
```

```
[6]: trainData, validData, testData, trainTarget, validTarget, testTarget = loadData()
print(f"Training Data: {trainData.shape}\tTraining tagets: {trainTarget.shape}")
print(f"Validation Data: {validData.shape}\tValidation tagets: {validTarget.
→shape}")
print(f"Testing Data: {testData.shape}\tTesting tagets:{testTarget.shape}")
```

Training Data: (15000, 28, 28) Training tagets: (15000,)  
Validation Data: (1000, 28, 28) Validation tagets: (1000,)  
Testing Data: (2724, 28, 28) Testing tagets:(2724,)

```
[0]: def plot(image, target, ax=None):
    ax = plt.gca() if ax == None else ax
    ax.imshow(image, cmap=plt.cm.gray)
    target_names = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J']
    ax.set_title(target_names[target])
    # targets interger encoded from 0 to 9 corresponding to 'A' to 'J',
→respectively
```

```
[8]: fig, axis = plt.subplots(2, 5, figsize=(16, 5))
for ax in axis.reshape(-1):
    r = np.random.randint(trainData.shape[0])
    plot(trainData[r], trainTarget[r], ax=ax)
plt.show()
```



## 0.2 Useful Functions

Some useful functions that will be used throughout the assignment such as getting random weights, getting the accuracy of a batch, making the loss and accuracy plots look nice, and global variables used throughout the code

```
[0]: # given by the assignment
def convertOneHot(trainTarget, validTarget, testTarget):
    newtrain = np.zeros((trainTarget.shape[0], 10))
    newvalid = np.zeros((validTarget.shape[0], 10))
    newtest = np.zeros((testTarget.shape[0], 10))
    for item in range(0, trainTarget.shape[0]):
        newtrain[item][trainTarget[item]] = 1
    for item in range(0, validTarget.shape[0]):
        newvalid[item][validTarget[item]] = 1
    for item in range(0, testTarget.shape[0]):
        newtest[item][testTarget[item]] = 1
    return newtrain, newvalid, newtest

[0]: def accuracy(y_pred, y):
    if y_pred.shape != y.shape:
        raise ValueError(f"prediction dimension {y_pred.shape} and label_
→dimensions {y.shape} don't match")
    return np.sum(y_pred.argmax(axis=1) == y.argmax(axis=1)) / y.shape[0]

[0]: def plot_loss(x, train_loss=None, valid_loss=None, test_loss=None, title=None,
→ax=None):
    ax = plt.gca() if ax == None else ax
    if train_loss != None:
        ax.plot(x, train_loss, label="Training Loss")
    if valid_loss != None:
        ax.plot(x, valid_loss, label="Validation Loss")
    if test_loss != None:
        ax.plot(x, test_loss, label="Testing Loss")

    ax.set_title("Loss" if title == None else title)

    ax.set_xlabel("Iterations")
    ax.set_xlim(left=0)
    ax.set_ylabel("Loss")
    ax.set_ylim(bottom=0)
    ax.legend(loc="upper right")

def plot_accuracy(x, train_accuracy=None, valid_accuracy=None,
→test_accuracy=None, title=None, ax=None):
    ax = plt.gca() if ax == None else ax
    if train_accuracy != None:
        ax.plot(x, train_accuracy, label="Training Accuracy")
```

```

if valid_accuracy != None:
    ax.plot(x, valid_accuracy, label="Validation Accuracy")
if test_accuracy != None:
    ax.plot(x, test_accuracy, label="Testing Accuracy")

ax.set_title("Accuracy" if title == None else title)

ax.set_xlabel("Iterations")
ax.set_xlim(left=0)
ax.set_ylabel("Accuracy")
ax.set_yticks(np.arange(0, 1.1, step=0.1))
ax.grid(linestyle='-', axis='y')
ax.legend(loc="lower right")

def display_statistics(train_loss=None, train_acc=None, valid_loss=None,
    →valid_acc=None,
                        test_loss=None, test_acc=None, num=True, plot=True):

    t1 = "-" if train_loss is None else round(train_loss[-1], 4)
    ta = "-" if train_acc is None else round(train_acc[-1]*100, 2)
    vl = "-\t" if valid_loss is None else round(valid_loss[-1], 4)
    va = "-" if valid_acc is None else round(valid_acc[-1]*100, 2)
    sl = "-\t\t" if test_loss is None else round(test_loss[-1], 4)
    sa = "-" if test_acc is None else round(test_acc[-1]*100, 2)

    if num:
        print(f"Training loss: {t1}{':.20s}\t\tTraining acc: {ta}{':' if ta != ' '
    →'- ' else ''}")
        print(f"Validation loss: {vl}{':.20s}\tValidation acc: {va}{':' if va !
    →= ' ' else ''}")
        print(f"Testing loss: {sl}{':.20s}\tTesting acc: {sa}{':' if sa != ' '
    →else ''}")

    if plot:
        fig, ax = plt.subplots(1, 2, figsize=(18, 6))
        plot_loss(np.arange(0, len(train_loss), 1), train_loss, valid_loss,
    →test_loss, ax=ax[0])
        plot_accuracy(np.arange(0, len(train_loss), 1), train_acc, valid_acc,
    →test_acc, ax=ax[1])
        plt.show()
        plt.close()

```

```

[0]: TINY = 1e-20
newtrain, newvalid, newtest = convertOneHot(trainTarget, validTarget, testTarget)

```

```
VTDatasets = {"validData" : validData.reshape(validData.shape[0], -1),  
→ "validTarget" : newvalid,  
              "testData" : testData.reshape(testData.shape[0], -1), "testTarget"  
→ : newtest}  
  
N = trainData.shape[0]  
d = trainData.shape[1] * trainData.shape[2]  
K = 10
```

# 1 Neural Networks using Numpy

## 1.1 Helper Functions

```
[0]: def relu(x):  
    return np.maximum(0, x)  
  
def softmax(x):  
    return np.exp(x) / np.exp(x).sum()  
  
def softmax_batch(X):  
    return np.exp(X) / np.exp(X).sum(axis=1, keepdims=True)  
  
[0]: def computeLayer(X, W, b):  
    return X @ W.T + b  
  
[0]: # target is one-hot encoded  
def averageCE(target, prediction):  
    return -(target * np.log(prediction+TINY)).sum(axis=1).mean()  
  
# target is one-hot encoded  
def gradCE(target, predication):  
    return predication - target
```

## 1.2 Backpropagation Derivation

**Derivative of Softmax**  $p_i = \text{softmax}(\mathbf{o})_i = \frac{e^{o_i}}{\sum_{k=1}^K e^{o_k}}$

if  $i \neq j$

$$\frac{\partial p_j}{\partial o_i} = \frac{0 \cdot \sum_{k=1}^K e^{o_k} - e^{o_i} \cdot e^{o_j}}{\left(\sum_{k=1}^K e^{o_k}\right)^2} = \boxed{-p_i \cdot p_j}$$

if  $i = j$

$$\frac{\partial p_j}{\partial o_i} = \frac{e^{o_i} \cdot \sum_{k=1}^K e^{o_k} - e^{o_i} \cdot e^{o_j}}{\left(\sum_{k=1}^K e^{o_k}\right)^2} = \boxed{(1 - p_j) \cdot p_i}$$

**Derivative of Softmax + Cross Entropy Loss**  $L_{CE}(\mathbf{y}, \mathbf{p}) = -\sum_{k=1}^K y_k \log p_k$

$$\frac{\partial L_{CE}}{\partial o_i} = -\sum_{k=1}^K \frac{y_k}{p_k} \cdot \frac{\partial p_k}{\partial o_i} = -y_i(1 - p_i) - \sum_{k \neq i} \frac{y_k}{p_k} \cdot (-p_k p_i) = -y_i + y_i p_i + \sum_{k \neq i} y_k p_i = -y_i + p_i \cdot \sum_{k=1}^K y_k = p_i - y_i$$

In Vector Form:  $\boxed{\frac{\partial L_{CE}}{\partial \mathbf{o}} = \mathbf{p} - \mathbf{y}}$



**Remaining Backpropagation**  $\mathbf{o} = W_o \mathbf{g} + \mathbf{b}_o$

$$\frac{\partial L}{\partial W_o} = \frac{\partial L}{\partial \mathbf{o}} \cdot \left( \frac{\partial \mathbf{o}}{\partial W_o} \right)^T = \frac{\partial L}{\partial \mathbf{o}} \cdot \mathbf{g}^T$$

$$\frac{\partial L}{\partial \mathbf{b}_o} = \frac{\partial L}{\partial \mathbf{o}} \cdot \left( \frac{\partial \mathbf{o}}{\partial \mathbf{b}_o} \right)^T = \frac{\partial L}{\partial \mathbf{o}}$$

$$g_i = \text{ReLU}(h_i) = \max(h_i, 0)$$

$$\frac{\partial L}{\partial h_i} = \frac{\partial L}{\partial g_i} \cdot \frac{\partial g_i}{\partial h_i} = \begin{cases} \frac{\partial L}{\partial g_i} & \text{if } h_i > 0 \\ 0 & \text{if } h_i < 0 \end{cases}$$

$$\mathbf{h} = W_h \mathbf{x} + \mathbf{b}_h$$

$$\frac{\partial L}{\partial W_h} = \frac{\partial L}{\partial \mathbf{h}} \cdot \left( \frac{\partial \mathbf{h}}{\partial W_h} \right)^T = \frac{\partial L}{\partial \mathbf{h}} \cdot \mathbf{x}^T$$

$$\frac{\partial L}{\partial \mathbf{b}_h} = \frac{\partial L}{\partial \mathbf{h}} \cdot \left( \frac{\partial \mathbf{h}}{\partial \mathbf{b}_h} \right)^T = \frac{\partial L}{\partial \mathbf{h}}$$

### 1.3 Learning

```
[16]: class mini_NN(object):

    """
    Network Structure:
        input: x
        hidden: h = W_h * x + b_h
                g = ReLU(h)
        output: o = W_o * g + b_o
                p = softmax(o)
    """

    def __init__(self, D, F, K):
        # D, F, and K are the number of neurons in the input, hidden, and output
        → layers
        self.D = D
        self.F = F
        self.K = K
        self.init_weights()

    def init_weights(self):
        # getting random parameters using Xavier initialization scheme
        self.W_h = np.random.normal(0, np.sqrt(2.0/(self.D+self.F)), (self.F,
        → self.D))
```

```

        self.b_h = np.random.normal(0, np.sqrt(2.0/(self.D+self.F)), self.F)
        self.W_o = np.random.normal(0, np.sqrt(2.0/(self.F+self.K)), (self.K,
→self.F))
        self.b_o = np.random.normal(0, np.sqrt(2.0/(self.F+self.K)), self.K)

    def feedforward(self, X):
        # python can dynamically create attributes
        self.H = computeLayer(X, self.W_h, self.b_h)
        self.G = relu(self.H)
        self.O = computeLayer(self.G, self.W_o, self.b_o)
        self.P = softmax_batch(self.O)
        return self.P

    def backpropagation(self, X, y):
        # This function assumes that feedforward was called before,
        # which instantiates the needed activations

        # output layer activations
        dL_do = gradCE(y, self.P)

        # output layer parameters
        dL_dWo = dL_do.T @ self.G
        dL_dbo = dL_do

        # hidden layer activations
        dL_dg = dL_do @ self.W_o
        dL_dh = dL_dg.copy()
        dL_dh[self.H <= 0] = 0

        # hidden layer parameters
        dL_dWh = dL_dh.T @ X
        dL_dbh = dL_dh

        return dL_dWo , dL_dbo.sum(axis=0), dL_dWh, dL_dbh.sum(axis=0)

    def train(self, X, y, epochs=200, gamma=0.99, alpha=1e-5, F=None,
              validData=None, validTarget=None, testData=None, testTarget=None):
        # initializations
        self.F = self.F if F is None else F
        self.init_weights()

        train_loss, train_acc = [], []
        valid_loss, valid_acc = [], []
        test_loss, test_acc = [], []

        v_Wo, v_Wh = 0, 0

```

```

    for e in range(epochs):
        if e > 0:
            print(f"epoch: {e+1}\tloss: {train_loss[-1]:.4f}\tacc:␣
→{train_acc[-1]:.4f}")
        else:
            print("epoch:", e+1)

        # getting predictions
        p = self.feedforward(X)
        train_loss.append( averageCE(p, y) )
        train_acc.append( accuracy(p, y) )

        # getting gradients
        dL_dWo, dL_dbo, dL_dWh, dL_dbh = self.backpropagation(X, y)

        # updating parameters
        v_Wo = gamma * v_Wo + alpha * dL_dWo
        self.W_o -= v_Wo

        self.b_o -= alpha * dL_dbo

        v_Wh = gamma * v_Wh + alpha * dL_dWh
        self.W_h -= v_Wh

        self.b_h -= alpha * dL_dbh

        # calculating statistics
        if not validData is None and not validTarget is None:
            p = self.feedforward(validData)
            valid_loss.append(averageCE(p, validTarget))
            valid_acc.append(accuracy(p, validTarget))
        if not testData is None and not testTarget is None:
            p = self.feedforward(testData)
            test_loss.append(averageCE(p, testTarget))
            test_acc.append(accuracy(p, testTarget))

        statistics = (train_loss, train_acc)
        if not validData is None and not validTarget is None:
            statistics += (valid_loss, valid_acc, )
        if not testData is None and not testTarget is None:
            statistics += (test_loss, test_acc,)
        return statistics

X = trainData.reshape(N, d)

y = newtrain

```

```

# For investigation, analyze how long each hyperparameter set takes to train
import time
start = time.time()

model = mini_NN(d, 1000, K)
statistics = model.train(X, y, epochs=200, gamma=0.99, alpha=1e-5, **VTDatasets)
display_statistics(*statistics)
print(f"Time is {time.time() - start}.")

```

```

epoch: 1
epoch: 2      loss: 41.8202  acc: 0.0623
epoch: 3      loss: 38.6344  acc: 0.1826
epoch: 4      loss: 30.5366  acc: 0.5814
epoch: 5      loss: 23.6037  acc: 0.6173
epoch: 6      loss: 20.7250  acc: 0.7061
epoch: 7      loss: 17.2101  acc: 0.7915
epoch: 8      loss: 19.0484  acc: 0.6860
epoch: 9      loss: 15.4555  acc: 0.7875
epoch: 10     loss: 13.4947  acc: 0.7948
epoch: 11     loss: 12.6952  acc: 0.7845
epoch: 12     loss: 12.3658  acc: 0.8017
epoch: 13     loss: 11.3749  acc: 0.8021
epoch: 14     loss: 9.8147   acc: 0.8311
epoch: 15     loss: 9.0415   acc: 0.8425
epoch: 16     loss: 8.4522   acc: 0.8576
epoch: 17     loss: 9.6073   acc: 0.8306
epoch: 18     loss: 8.6273   acc: 0.8579
epoch: 19     loss: 7.4822   acc: 0.8685
epoch: 20     loss: 7.6507   acc: 0.8607
epoch: 21     loss: 7.6104   acc: 0.8610
epoch: 22     loss: 7.3691   acc: 0.8727
epoch: 23     loss: 7.8118   acc: 0.8719
epoch: 24     loss: 7.4743   acc: 0.8791
epoch: 25     loss: 7.5346   acc: 0.8783
epoch: 26     loss: 7.4533   acc: 0.8836
epoch: 27     loss: 7.5341   acc: 0.8862
epoch: 28     loss: 7.8440   acc: 0.8823
epoch: 29     loss: 7.6133   acc: 0.8800
epoch: 30     loss: 7.1122   acc: 0.8833
epoch: 31     loss: 6.7606   acc: 0.8868
epoch: 32     loss: 6.5217   acc: 0.8953
epoch: 33     loss: 6.7583   acc: 0.8958
epoch: 34     loss: 6.7986   acc: 0.8963
epoch: 35     loss: 6.5820   acc: 0.8950
epoch: 36     loss: 6.3596   acc: 0.8940
epoch: 37     loss: 6.0831   acc: 0.8988

```

epoch: 38	loss: 5.8347	acc: 0.9041
epoch: 39	loss: 5.7389	acc: 0.9072
epoch: 40	loss: 5.7982	acc: 0.9054
epoch: 41	loss: 5.8156	acc: 0.9033
epoch: 42	loss: 5.7793	acc: 0.9019
epoch: 43	loss: 5.7104	acc: 0.9035
epoch: 44	loss: 5.5470	acc: 0.9058
epoch: 45	loss: 5.3071	acc: 0.9114
epoch: 46	loss: 5.1573	acc: 0.9159
epoch: 47	loss: 5.1544	acc: 0.9161
epoch: 48	loss: 5.1194	acc: 0.9167
epoch: 49	loss: 5.0175	acc: 0.9161
epoch: 50	loss: 5.0072	acc: 0.9149
epoch: 51	loss: 5.0473	acc: 0.9134
epoch: 52	loss: 5.0269	acc: 0.9153
epoch: 53	loss: 4.9502	acc: 0.9182
epoch: 54	loss: 4.8594	acc: 0.9213
epoch: 55	loss: 4.7863	acc: 0.9250
epoch: 56	loss: 4.7844	acc: 0.9261
epoch: 57	loss: 4.8328	acc: 0.9245
epoch: 58	loss: 4.8698	acc: 0.9230
epoch: 59	loss: 4.8308	acc: 0.9230
epoch: 60	loss: 4.7091	acc: 0.9255
epoch: 61	loss: 4.5888	acc: 0.9278
epoch: 62	loss: 4.5251	acc: 0.9293
epoch: 63	loss: 4.4795	acc: 0.9290
epoch: 64	loss: 4.4006	acc: 0.9298
epoch: 65	loss: 4.3017	acc: 0.9309
epoch: 66	loss: 4.2276	acc: 0.9315
epoch: 67	loss: 4.1908	acc: 0.9337
epoch: 68	loss: 4.1885	acc: 0.9333
epoch: 69	loss: 4.1939	acc: 0.9337
epoch: 70	loss: 4.1573	acc: 0.9341
epoch: 71	loss: 4.0575	acc: 0.9361
epoch: 72	loss: 3.9474	acc: 0.9387
epoch: 73	loss: 3.8890	acc: 0.9396
epoch: 74	loss: 3.8850	acc: 0.9396
epoch: 75	loss: 3.8930	acc: 0.9409
epoch: 76	loss: 3.8705	acc: 0.9416
epoch: 77	loss: 3.8124	acc: 0.9427
epoch: 78	loss: 3.7506	acc: 0.9445
epoch: 79	loss: 3.7120	acc: 0.9455
epoch: 80	loss: 3.6882	acc: 0.9465
epoch: 81	loss: 3.6495	acc: 0.9467
epoch: 82	loss: 3.5668	acc: 0.9474
epoch: 83	loss: 3.4519	acc: 0.9483
epoch: 84	loss: 3.3431	acc: 0.9499
epoch: 85	loss: 3.2696	acc: 0.9510


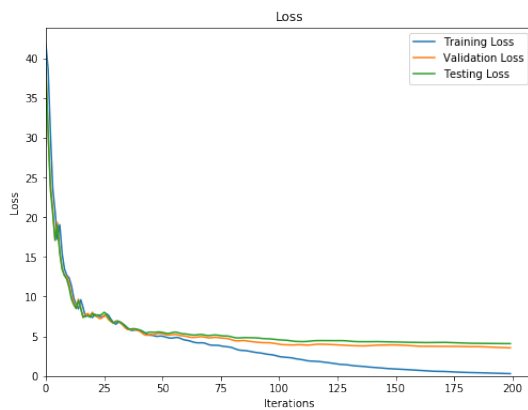
epoch: 86	loss: 3.2365	acc: 0.9515
epoch: 87	loss: 3.2230	acc: 0.9508
epoch: 88	loss: 3.1943	acc: 0.9508
epoch: 89	loss: 3.1376	acc: 0.9517
epoch: 90	loss: 3.0705	acc: 0.9536
epoch: 91	loss: 3.0146	acc: 0.9551
epoch: 92	loss: 2.9787	acc: 0.9553
epoch: 93	loss: 2.9516	acc: 0.9559
epoch: 94	loss: 2.9135	acc: 0.9561
epoch: 95	loss: 2.8565	acc: 0.9574
epoch: 96	loss: 2.7934	acc: 0.9586
epoch: 97	loss: 2.7478	acc: 0.9593
epoch: 98	loss: 2.7193	acc: 0.9608
epoch: 99	loss: 2.6844	acc: 0.9615
epoch: 100	loss: 2.6219	acc: 0.9625
epoch: 101	loss: 2.5404	acc: 0.9639
epoch: 102	loss: 2.4663	acc: 0.9646
epoch: 103	loss: 2.4184	acc: 0.9649
epoch: 104	loss: 2.3951	acc: 0.9649
epoch: 105	loss: 2.3812	acc: 0.9657
epoch: 106	loss: 2.3583	acc: 0.9661
epoch: 107	loss: 2.3147	acc: 0.9663
epoch: 108	loss: 2.2548	acc: 0.9680
epoch: 109	loss: 2.1970	acc: 0.9692
epoch: 110	loss: 2.1519	acc: 0.9697
epoch: 111	loss: 2.1104	acc: 0.9705
epoch: 112	loss: 2.0592	acc: 0.9714
epoch: 113	loss: 1.9977	acc: 0.9730
epoch: 114	loss: 1.9409	acc: 0.9739
epoch: 115	loss: 1.9038	acc: 0.9751
epoch: 116	loss: 1.8881	acc: 0.9751
epoch: 117	loss: 1.8827	acc: 0.9745
epoch: 118	loss: 1.8700	acc: 0.9745
epoch: 119	loss: 1.8424	acc: 0.9747
epoch: 120	loss: 1.8024	acc: 0.9766
epoch: 121	loss: 1.7576	acc: 0.9779
epoch: 122	loss: 1.7173	acc: 0.9786
epoch: 123	loss: 1.6858	acc: 0.9793
epoch: 124	loss: 1.6565	acc: 0.9799
epoch: 125	loss: 1.6186	acc: 0.9804
epoch: 126	loss: 1.5700	acc: 0.9811
epoch: 127	loss: 1.5217	acc: 0.9816
epoch: 128	loss: 1.4884	acc: 0.9823
epoch: 129	loss: 1.4716	acc: 0.9827
epoch: 130	loss: 1.4576	acc: 0.9827
epoch: 131	loss: 1.4301	acc: 0.9832
epoch: 132	loss: 1.3875	acc: 0.9833
epoch: 133	loss: 1.3418	acc: 0.9840

epoch: 134	loss: 1.3063	acc: 0.9853
epoch: 135	loss: 1.2854	acc: 0.9853
epoch: 136	loss: 1.2719	acc: 0.9855
epoch: 137	loss: 1.2554	acc: 0.9853
epoch: 138	loss: 1.2300	acc: 0.9858
epoch: 139	loss: 1.1959	acc: 0.9864
epoch: 140	loss: 1.1591	acc: 0.9870
epoch: 141	loss: 1.1282	acc: 0.9873
epoch: 142	loss: 1.1075	acc: 0.9877
epoch: 143	loss: 1.0932	acc: 0.9881
epoch: 144	loss: 1.0777	acc: 0.9885
epoch: 145	loss: 1.0536	acc: 0.9891
epoch: 146	loss: 1.0210	acc: 0.9899
epoch: 147	loss: 0.9879	acc: 0.9897
epoch: 148	loss: 0.9628	acc: 0.9903
epoch: 149	loss: 0.9474	acc: 0.9903
epoch: 150	loss: 0.9357	acc: 0.9902
epoch: 151	loss: 0.9207	acc: 0.9904
epoch: 152	loss: 0.9005	acc: 0.9904
epoch: 153	loss: 0.8784	acc: 0.9908
epoch: 154	loss: 0.8589	acc: 0.9912
epoch: 155	loss: 0.8433	acc: 0.9913
epoch: 156	loss: 0.8296	acc: 0.9920
epoch: 157	loss: 0.8149	acc: 0.9926
epoch: 158	loss: 0.7980	acc: 0.9926
epoch: 159	loss: 0.7787	acc: 0.9927
epoch: 160	loss: 0.7587	acc: 0.9927
epoch: 161	loss: 0.7393	acc: 0.9929
epoch: 162	loss: 0.7216	acc: 0.9932
epoch: 163	loss: 0.7062	acc: 0.9931
epoch: 164	loss: 0.6915	acc: 0.9931
epoch: 165	loss: 0.6746	acc: 0.9933
epoch: 166	loss: 0.6547	acc: 0.9939
epoch: 167	loss: 0.6337	acc: 0.9939
epoch: 168	loss: 0.6162	acc: 0.9941
epoch: 169	loss: 0.6050	acc: 0.9945
epoch: 170	loss: 0.6002	acc: 0.9948
epoch: 171	loss: 0.5981	acc: 0.9946
epoch: 172	loss: 0.5938	acc: 0.9946
epoch: 173	loss: 0.5833	acc: 0.9948
epoch: 174	loss: 0.5662	acc: 0.9949
epoch: 175	loss: 0.5455	acc: 0.9952
epoch: 176	loss: 0.5258	acc: 0.9955
epoch: 177	loss: 0.5108	acc: 0.9956
epoch: 178	loss: 0.5011	acc: 0.9958
epoch: 179	loss: 0.4948	acc: 0.9961
epoch: 180	loss: 0.4887	acc: 0.9961
epoch: 181	loss: 0.4800	acc: 0.9960

```

epoch: 182      loss: 0.4688      acc: 0.9965
epoch: 183      loss: 0.4572      acc: 0.9965
epoch: 184      loss: 0.4473      acc: 0.9965
epoch: 185      loss: 0.4399      acc: 0.9965
epoch: 186      loss: 0.4343      acc: 0.9963
epoch: 187      loss: 0.4288      acc: 0.9963
epoch: 188      loss: 0.4220      acc: 0.9967
epoch: 189      loss: 0.4134      acc: 0.9965
epoch: 190      loss: 0.4034      acc: 0.9965
epoch: 191      loss: 0.3930      acc: 0.9965
epoch: 192      loss: 0.3827      acc: 0.9967
epoch: 193      loss: 0.3728      acc: 0.9969
epoch: 194      loss: 0.3631      acc: 0.9971
epoch: 195      loss: 0.3540      acc: 0.9973
epoch: 196      loss: 0.3460      acc: 0.9974
epoch: 197      loss: 0.3394      acc: 0.9978
epoch: 198      loss: 0.3334      acc: 0.9978
epoch: 199      loss: 0.3273      acc: 0.9977
epoch: 200      loss: 0.3204      acc: 0.9981
Training loss: 0.3129      Training acc: 99.81%
Validation loss: 3.558      Validation acc: 92.4%
Testing loss: 4.0922      Testing acc: 91.3%

```



Time is 342.4068307876587.

## 1.4 Hyperparameter Investigation

### 0.1 Number of Hidden Units

#### 0.1.1 Number of hidden units: 100

Statistic	Value	Statistic	Value
Training loss:	4.6824	Training acc:	92.31%
Validation loss:	5.5578	Validation acc:	89.8%



Statistic	Value	Statistic	Value
Testing loss:	5.7423	Testing acc:	89.46%
Training Time(s):	42.705		

**Comments:** Accuracy seems to have converged by 50 iterations. Training this network with 100 hidden units was the fastest of the three. This is likely because there were very few parameters to optimize.

### 0.1.2 Number of hidden units: 500

Statistic	Value	Statistic	Value
Training loss:	0.7009	Training acc:	99.36%
Validation loss:	3.894	Validation acc:	91.9%
Testing loss:	4.1808	Testing acc:	91.19%
Training Time(s):	176.361		

**Comments:** Accuracy seems to have converged by 50 iterations. Training this network with 500 hidden units was slower than with 100 units, but the validation/testing accuracies showed slight improvement. Further, Training loss and training accuracy near perfection. Since this trend is not matched by validation/testing accuracies, we see that the network is overfitting as it is learning the training examples too well.

### 0.1.3 Number of hidden units: 2000

Statistic	Value	Statistic	Value
Training loss:	0.2303	Training acc:	99.82%
Validation loss:	3.3766	Validation acc:	93.0%
Testing loss:	3.8814	Testing acc:	91.85%
Training Time(s):	645.494		

**Comments:** Accuracy seems to have converged by 50 iterations. Training this network with 2000 hidden units was considerably slower than both other networks, despite nearly identical validation/testing accuracies. Training loss and training accuracy are essentially perfect by the end of 200 epochs. Since this trend is not matched by validation/testing accuracies, we see that the network is overfitting considerably. It has essentially memorized the training examples, but can not generalize as well (beyond the performance attained by smaller networks) to validation/testing data.

**General comments:** Also, the small bump and increase in loss before once again decreasing suggest that the parameters went through a local minimum before converging. Momentum likely helped it converge faster as it doesn't appear to have gotten stuck in the local min. The smaller networks trained considerably faster, but had lower accuracy (even though the difference was very small). The larger networks clearly overfit the training data; this indicates that they are too complex for their classification tasks and/or the input data should be improved.

## 0.2 Early Stopping

From the plots, we observe that training should have stopped at an early stopping point of 50 iterations. Beyond this point, there is very little improvement in either validation/testing losses or accuracies.

## 2 Neural Networks in Tensorflow

### 2.1 Model implementation

```
[0]: # load + reshape data
trainData, validData, testData, trainTarget, validTarget, testTarget = loadData()
trainData = trainData.reshape(15000,28,28,1)
validData = validData.reshape(1000,28,28,1)
testData = testData.reshape(2724,28,28,1)

# one-hot encode
train_labels, valid_labels, test_labels = convertOneHot(trainTarget,
    →validTarget, testTarget)

# training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# create model
model = models.Sequential()
model.add(layers.InputLayer(input_shape=(28, 28,1)))
    →# input layer
model.add(layers.Conv2D(
    →# conv layer
        filters=32,
        strides=(1,1),
        kernel_size=[3, 3],
        padding="same",
        activation='relu',
        kernel_initializer=tf.contrib.layers.xavier_initializer(uniform=False)))
model.add(layers.BatchNormalization())
    →# batch norm
model.add(layers.MaxPooling2D((2, 2)))
    →# max pooling
model.add(layers.Flatten())
    →# flatten
model.add(layers.Dense(784, activation='relu'))
    →# fully-connected 784 w/ ReLu
model.add(layers.Dense(10))
    →# fully-connected 10
model.add(layers.Softmax())
    →# softmax output

# compile model w/ Adam optimizer + cross entropy loss
model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
               loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
```

```
metrics=['accuracy'])
```

## 2.2 Model Training

```
[18]: # callback to test after each epoch
class TestCallback(tf.keras.callbacks.Callback):
    def __init__(self, test_data):
        self.test_data = test_data
        self.test_acc = []
        self.test_loss = []

    def on_epoch_end(self, epoch, logs=None):
        # perform a test per epoch
        x, y = self.test_data
        loss, acc = self.model.evaluate(x, y, verbose=0, batch_size=32)
        self.test_loss.append(loss)
        self.test_acc.append(acc)
        # append to returned dictionary
        logs["test_loss"] = self.test_loss
        logs["test_acc"] = self.test_acc

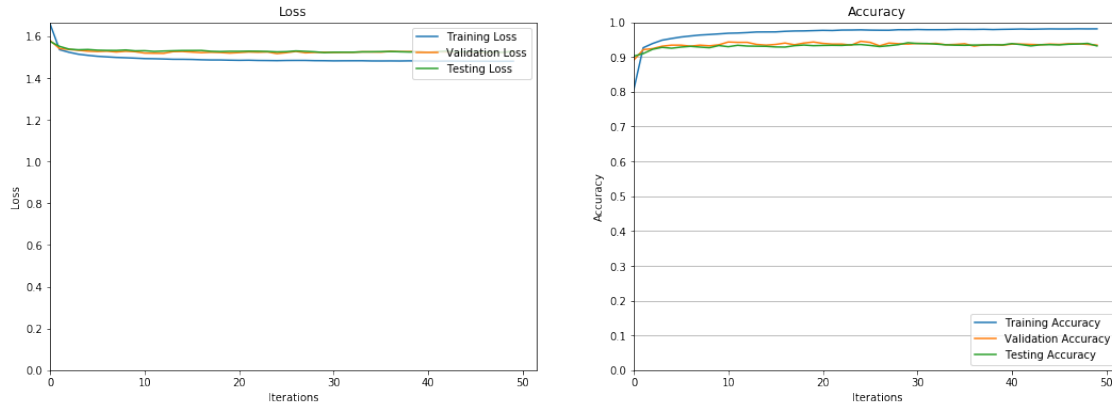
# train
history = model.fit(trainData, train_labels,
                    validation_data=(validData, valid_labels),
                    epochs=epochs,
                    batch_size=batch_size,
                    callbacks=[TestCallback((testData, test_labels))],
                    verbose=0, # 0 = silent, 1 = per epoch
                    shuffle=True)

# print(history.history)

# plot accuracy + loss
train_acc = history.history["acc"]
val_acc = history.history["val_acc"]
test_acc = history.history["test_acc"][0]
train_loss = history.history["loss"]
val_loss = history.history["val_loss"]
test_loss = history.history["test_loss"][0]

statistics = (train_loss, train_acc, val_loss, val_acc, test_loss, test_acc)
display_statistics(*statistics)
```

```
Training loss: 1.4806           Training acc: 98.08%
Validation loss: 1.5268 Validation acc: 93.4%
Testing loss: 1.5286    Testing acc: 93.17%
```



## 2.3 Hyperparameter Investigation

[20]: # 2.3.1: L2 Regularization

```
# training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# test all weight decays [0.01, 0.1, 0.5]
for scale in [0.01, 0.1, 0.5]:
    print("\nL2 Normalization with {}".format(scale))

    # create model
    model = models.Sequential()
    model.add(layers.InputLayer(input_shape=(28, 28,1))) # input layer
    model.add(layers.Conv2D(
        # conv layer
        filters=32,
        strides=(1,1),
        kernel_size=[3, 3],
        padding="same",
        activation='relu',
        kernel_initializer=tf.contrib.layers.
        xavier_initializer(uniform=False)))
    model.add(layers.BatchNormalization())
    # batch norm
    model.add(layers.MaxPooling2D((2, 2)))
    # max pooling
    model.add(layers.Flatten())
    # flatten
    model.add(layers.Dense(784,
```

```

        activation='relu',
        kernel_regularizer=tf.contrib.layers.
→l2_regularizer(scale=scale))) # fully-connected 784 w/ ReLu
        model.add(layers.Dense(10))
→ # fully-connected 10
        model.add(layers.Softmax())
→ # softmax output

# compile model w/ Adam optimizer + cross entropy loss
model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
               loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
               metrics=['accuracy'])

# train
history_1 = model.fit(trainData, train_labels,
                      validation_data = (validData, valid_labels),
                      epochs=epochs,
                      batch_size=batch_size,
                      callbacks=[TestCallback((testData, test_labels))],
                      verbose=0, # 0 = silent, 1 = per epoch
                      shuffle=True)

# display stats
train_acc = history_1.history["acc"]
val_acc = history_1.history["val_acc"]
test_acc = history_1.history["test_acc"][0]
train_loss = history_1.history["loss"]
val_loss = history_1.history["val_loss"]
test_loss = history_1.history["test_loss"][0]

statistics = (train_loss, train_acc, val_loss, val_acc, test_loss, test_acc)
display_statistics(*statistics)

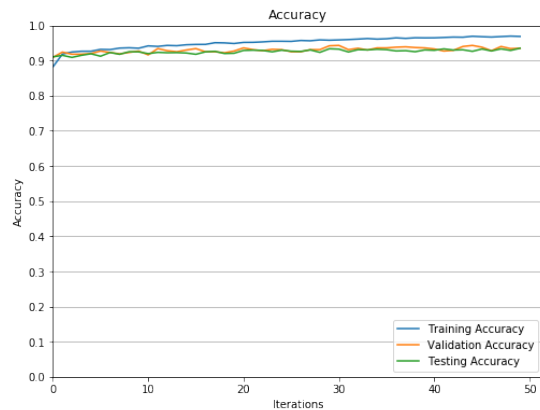
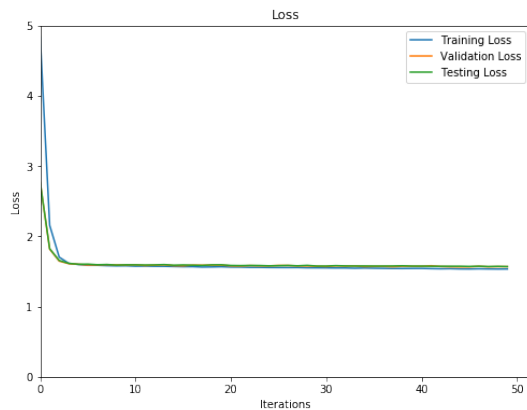
```

L2 Normalization with 0.01

```

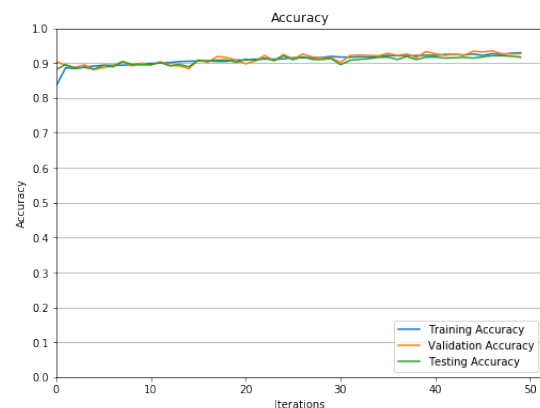
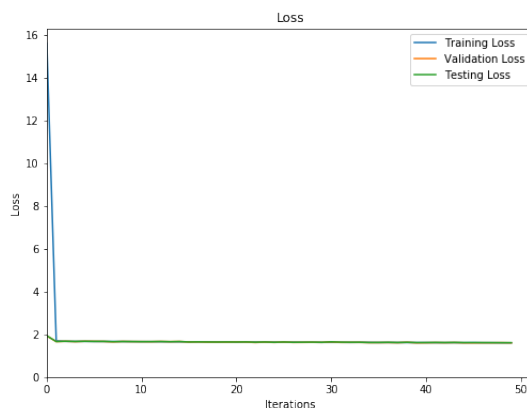
Training loss: 1.5345           Training acc: 96.86%
Validation loss: 1.5678 Validation acc: 93.5%
Testing loss: 1.569           Testing acc: 93.5%

```



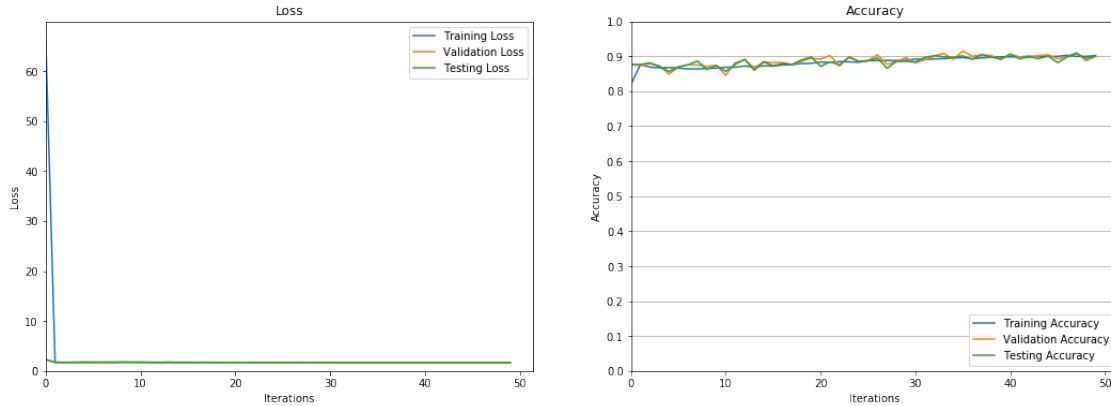
L2 Normalization with 0.1

Training loss: 1.6037                      Training acc: 92.97%  
 Validation loss: 1.6036      Validation acc: 92.7%  
 Testing loss: 1.6122              Testing acc: 91.7%



L2 Normalization with 0.5

Training loss: 1.6725                      Training acc: 90.21%  
 Validation loss: 1.6759      Validation acc: 90.0%  
 Testing loss: 1.6763              Testing acc: 90.01%



**General Comments:** L2 regularization is a technique to reduce overfitting. For small values of  $\lambda$ , we expect a slight improvement in model performance and a reduced discrepancy between training and validation/test accuracies. This is exactly what we observe here. For  $\lambda = 0.01$ , we observe validation and testing accuracy both improve slightly while training accuracy decreases by  $\sim 1.5\%$ . So, the model performance improves slightly and the model overfits less as expected. For values of  $\lambda$  that are too large, we expect the model to start underfitting as we are harshly penalizing parameter vectors which are large in magnitude. As  $\lambda$  increases, we see the discrepancy between training and validation/test accuracies decrease dramatically (the model is no longer overfitting). However, this change is accompanied by a reduction in overall accuracy (by  $\sim 2\%$ ) as the model is now underfitting.

```
[21]: # 2.3.2: Dropout

# training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# for rate in [0.9, 0.75, 0.5]:
for rate in [0.1, 0.25, 0.5]:
    print("\nDropout with probability {printed: .2f}\n".format(printed=1-rate))

    # create model
    model = models.Sequential()
    model.add(layers.InputLayer(input_shape=(28, 28,1))) # input layer
    model.add(layers.Conv2D(
        # conv layer
        filters=32,
        strides=(1,1),
        kernel_size=[3, 3],
        padding="same",
        activation='relu',
```



```

        kernel_initializer=tf.contrib.layers.
→xavier_initializer(uniform=False)))
        model.add(layers.BatchNormalization())
→    # batch norm
        model.add(layers.MaxPooling2D((2, 2)))
→    # max pooling
        model.add(layers.Flatten())
→    # flatten
        model.add(layers.Dense(784))
→    # fully-connected 784 w/ ReLu
        model.add(layers.Dropout(rate=rate))
→    # dropout
        model.add(layers.ReLU())
→    # Relu activation
        model.add(layers.Dense(10))
→    # fully-connected 10
        model.add(layers.Softmax())
→    # softmax output

# compile model w/ Adam optimizer + cross entropy loss
model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
               loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
               metrics=['accuracy'])

# train
history_2 = model.fit(trainData, train_labels,
                      validation_data = (validData, valid_labels),
                      epochs=epochs,
                      batch_size=batch_size,
                      callbacks=[TestCallback((testData, test_labels))],
                      verbose=0, # 0 = silent, 1 = per epoch
                      shuffle=True)

# display stats
train_acc = history_2.history["acc"]
val_acc = history_2.history["val_acc"]
test_acc = history_2.history["test_acc"][0]
train_loss = history_2.history["loss"]
val_loss = history_2.history["val_loss"]
test_loss = history_2.history["test_loss"][0]

statistics = (train_loss, train_acc, val_loss, val_acc, test_loss, test_acc)
display_statistics(*statistics)

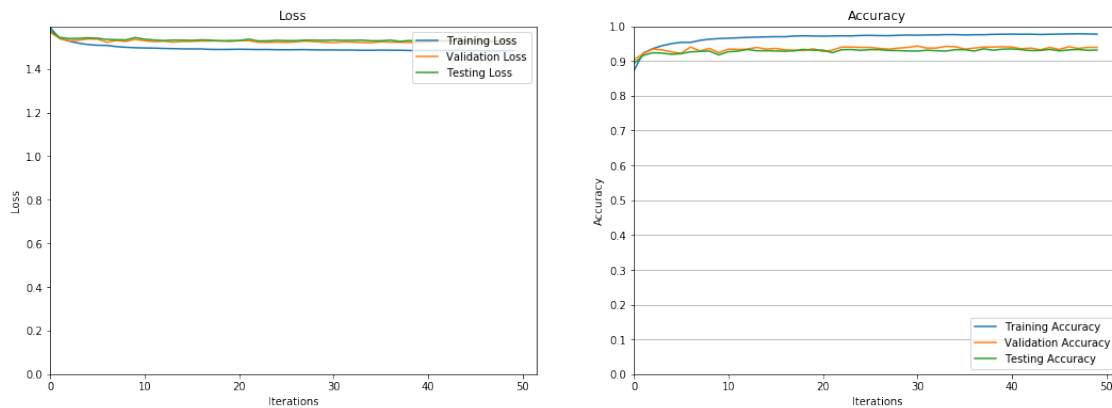
```

Dropout with probability 0.90

Training loss: 1.484

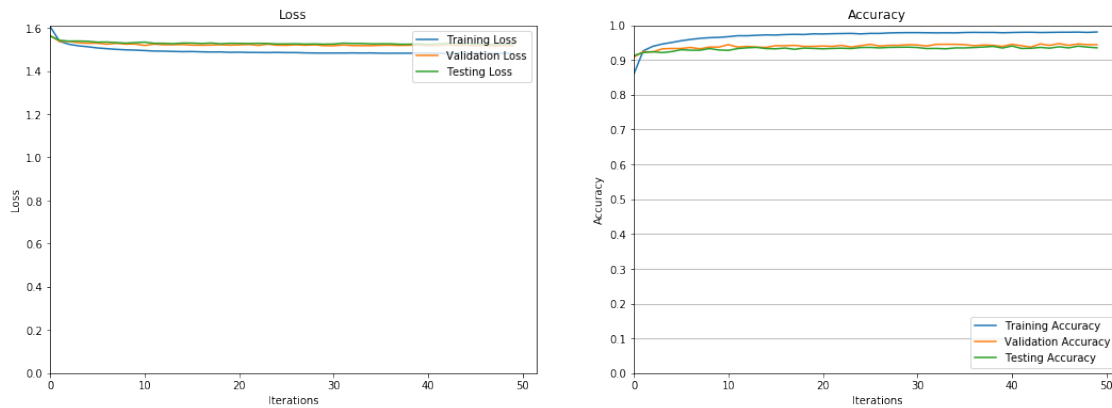
Training acc: 97.71%

Validation loss: 1.5219 Validation acc: 93.9%  
Testing loss: 1.5291 Testing acc: 93.14%



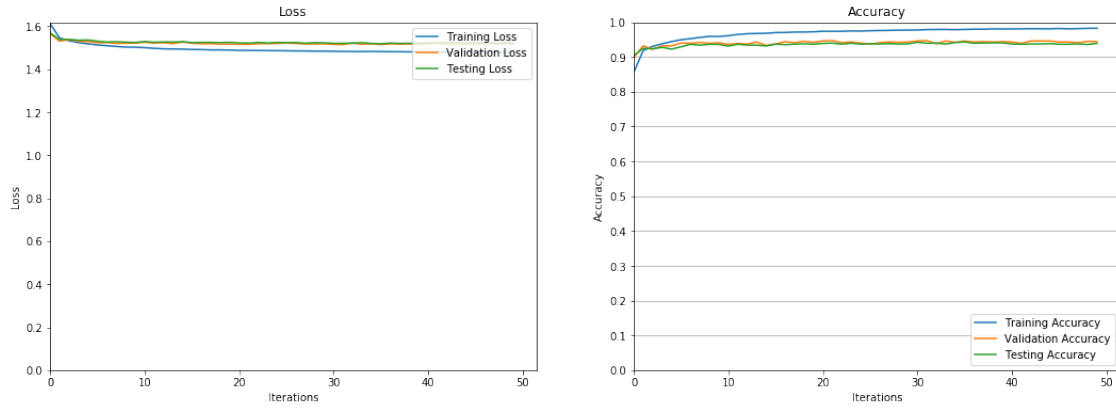
Dropout with probability 0.75

Training loss: 1.4808 Training acc: 98.04%  
Validation loss: 1.5174 Validation acc: 94.4%  
Testing loss: 1.5249 Testing acc: 93.5%



Dropout with probability 0.50

Training loss: 1.4788 Training acc: 98.25%  
Validation loss: 1.5178 Validation acc: 94.4%  
Testing loss: 1.523 Testing acc: 93.91%



**General Comments:** Dropout is a technique which aims to decrease overfitting in our model. We expect see this in the form of improved accuracy and loss statistics. Indeed, this is exactly observed. As the amount of dropout increases (keeping the rate within reason), the validation and testing accuracies both improve. Note that we also see training accuracy improve. So, it seems that the model as an increased capacity to learn overall. With dropout, it is able to learn the training data better, but it is also able to generalize better.