

A2WriteUp_part2

February 29, 2020

0 Setup

```
[0]: # ignore all future warnings
from warnings import simplefilter
simplefilter(action='ignore', category=FutureWarning)
```

```
[0]: # importing tensorflow
try:
    import google.colab
    import tensorflow as tf
    %tensorflow_version 1.13
except:
    import tensorflow as tf
    assert tf.__version__ == "1.13.1"

    # ignore tensorflow depreciation warnings
    import tensorflow.python.util.deprecation as deprecation
    deprecation._PRINT_DEPRECATION_WARNINGS = False
```

```
[0]: # imports
import numpy as np
import matplotlib.pyplot as plt
from tensorflow.keras import layers, models
```

```
[4]: print(tf.__version__)
```

1.15.0

0.1 Visualizing the Dataset

```
[0]: # given by the assignment
def loadData():
    with np.load("notMNIST.npz") as data:
        Data, Target = data["images"], data["labels"]
        np.random.seed(521)
        randIndx = np.arange(len(Data))
        np.random.shuffle(randIndx)
        Data = Data[randIndx]/255.
        Target = Target[randIndx]
        trainData, trainTarget = Data[:15000], Target[:15000]
        validData, validTarget = Data[15000:16000], Target[15000:16000]
        testData, testTarget = Data[16000:], Target[16000:]
    return trainData, validData, testData, trainTarget, validTarget, testTarget
```

```
[6]: trainData, validData, testData, trainTarget, validTarget, testTarget = loadData()
print(f"Training Data: {trainData.shape}\tTraining targets: {trainTarget.shape}")
print(f"Validation Data: {validData.shape}\tValidation targets: {validTarget.shape}")
print(f"Testing Data: {testData.shape}\tTesting targets:{testTarget.shape}")
```

```

FileNotFoundError                                Traceback (most recent call
last)

<ipython-input-6-bd37811a1dc2> in <module>()
----> 1 trainData, validData, testData, trainTarget, validTarget, testTarget
= loadData()
      2 print(f"Training Data: {trainData.shape}\tTraining targets:
{trainTarget.shape}")
      3 print(f"Validation Data: {validData.shape}\tValidation targets:
{validTarget.shape}")
      4 print(f"Testing Data: {testData.shape}\tTesting targets:{testTarget.
shape}")

<ipython-input-5-8c04b2024273> in loadData()
      1 def loadData():
----> 2     with np.load("notMNIST.npz") as data:
      3         Data, Target = data ["images"], data["labels"]
      4         np.random.seed(521)
      5         randIndx = np.arange(len(Data))

/usr/local/lib/python3.6/dist-packages/numpy/lib/npio.py in load(file,
mmap_mode, allow_pickle, fix_imports, encoding)
    426     own_fid = False
    427     else:
--> 428         fid = open(os_fspath(file), "rb")
    429         own_fid = True
    430
```

```
FileNotFoundError: [Errno 2] No such file or directory: 'notMNIST.npz'
```

```
[0]: def plot(image, target, ax=None):
      ax = plt.gca() if ax == None else ax
      ax.imshow(image, cmap=plt.cm.gray)
      target_names = ['A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J']
      ax.set_title(target_names[target])
      # targets interger encoded from 0 to 9 corresponding to 'A' to 'J',
      ↪ respectively
```

```
[0]: fig, axis = plt.subplots(2, 5, figsize=(16, 5))
      for ax in axis.reshape(-1):
          r = np.random.randint(trainData.shape[0])
          plot(trainData[r], trainTarget[r], ax=ax)
      plt.show()
```

0.2 Useful Functions

Some useful functions that will be used throughout the assignment such as getting random weights, getting the accuracy of a batch, making the loss and accuracy plots look nice, and global variables used throughout the code

```
[0]: # given by the assignment
def convertOneHot(trainTarget, validTarget, testTarget):
    newtrain = np.zeros((trainTarget.shape[0], 10))
    newvalid = np.zeros((validTarget.shape[0], 10))
    newtest = np.zeros((testTarget.shape[0], 10))
    for item in range(0, trainTarget.shape[0]):
        newtrain[item][trainTarget[item]] = 1
    for item in range(0, validTarget.shape[0]):
        newvalid[item][validTarget[item]] = 1
    for item in range(0, testTarget.shape[0]):
        newtest[item][testTarget[item]] = 1
    return newtrain, newvalid, newtest

[0]: def accuracy(y_pred, y):
    if y_pred.shape != y.shape:
        raise ValueError(f"prediction dimension {y_pred.shape} and label_
↳ dimensions {y.shape} don't match")
    return np.sum(y_pred.argmax(axis=1) == y.argmax(axis=1)) / y.shape[0]

[0]: def plot_loss(x, train_loss=None, valid_loss=None, test_loss=None, title=None,
↳ ax=None):
    ax = plt.gca() if ax == None else ax
    if train_loss != None:
        ax.plot(x, train_loss, label="Training Loss")
    if valid_loss != None:
        ax.plot(x, valid_loss, label="Validation Loss")
    if test_loss != None:
        ax.plot(x, test_loss, label="Testing Loss")

    ax.set_title("Loss" if title == None else title)

    ax.set_xlabel("Iterations")
    ax.set_xlim(left=0)
    ax.set_ylabel("Loss")
    ax.set_ylim(bottom=0)
    ax.legend(loc="upper right")

def plot_accuracy(x, train_accuracy=None, valid_accuracy=None,
↳ test_accuracy=None, title=None, ax=None):
    ax = plt.gca() if ax == None else ax
    if train_accuracy != None:
```



```
N = trainData.shape[0]
d = trainData.shape[1] * trainData.shape[2]
K = 10
```


2 Neural Networks in Tensorflow

2.1 Model implementation

```
[0]: # load + reshape data
trainData, validData, testData, trainTarget, validTarget, testTarget =  loadData()
trainData = trainData.reshape(15000,28,28,1)
validData = validData.reshape(1000,28,28,1)
testData = testData.reshape(2724,28,28,1)

# one-hot encode
train_labels, valid_labels, test_labels = convertOneHot(trainTarget,  validTarget, testTarget)

[0]: # training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# create model
model = models.Sequential()
model.add(layers.InputLayer(input_shape=(28, 28,1))) # input layer
model.add(layers.Conv2D( # conv layer
    filters=32,
    strides=(1,1),
    kernel_size=[3, 3],
    padding="same",
    activation='relu',
    kernel_initializer=tf.contrib.layers.xavier_initializer(uniform=False)))
model.add(layers.BatchNormalization()) # batch norm
model.add(layers.MaxPooling2D((2, 2))) # max pooling
model.add(layers.Flatten()) # flatten
model.add(layers.Dense(784, activation='relu')) # fully-connected 784
model.add(layers.Dense(10)) # fully-connected 10
model.add(layers.Softmax()) # softmax output

# compile model w/ Adam optimizer + cross entropy loss
model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
              loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
```

2.2 Model Training

```
[ ]: # callback to test after each epoch
class TestCallback(tf.keras.callbacks.Callback):
    def __init__(self, test_data):
        self.test_data = test_data
        self.test_acc = []
        self.test_loss = []



    def on_epoch_end(self, epoch, logs=None):
        # perform a test per epoch
        x, y = self.test_data
        loss, acc = self.model.evaluate(x, y, verbose=0, batch_size=32)
        self.test_loss.append(loss)
        self.test_acc.append(acc)
        # append to returned dictionary
        logs["test_loss"] = self.test_loss
        logs["test_acc"] = self.test_acc
```

```
[0]: # training
history = model.fit(trainData, train_labels,
                    validation_data = (validData, valid_labels),
                    epochs=epochs,
                    batch_size=batch_size,
                    callbacks=[TestCallback((testData, test_labels))],
                    verbose=0, # 0 = silent, 1 = per epoch
                    shuffle=True)

# display statistics
train_acc, train_loss = history.history["acc"], history.history["loss"]
val_acc, val_loss = history.history["val_acc"], history.history["val_loss"]
test_acc, test_loss = history.history["test_acc"][0], history.
    ↪history["test_loss"][0]

display_statistics(train_loss=train_loss, train_acc=train_acc,
                  valid_loss=val_loss, valid_acc=val_acc,
                  test_loss=test_loss, test_acc=test_acc)
```

Training loss: 1.4814 Training acc: 97.97%
Validation loss: 1.5248 Validation acc: 93.5%
Testing loss: 1.5293 Testing acc: 93.06%



2.3 Hyperparameter Investigation

2.3.1 L2 Regularization

```
[0]: # training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# test all weight decays [0.01, 0.1, 0.5]
for scale in [0.01, 0.1, 0.5]:
    print("\nL2 Normalization with {}".format(scale))

    # create model
    model = models.Sequential()
    model.add(layers.InputLayer(input_shape=(28, 28,1))) # input layer
    model.add(layers.Conv2D( # conv layer
        filters=32,
        strides=(1,1),
        kernel_size=[3, 3],
        padding="same",
        activation='relu',
        kernel_initializer=tf.contrib.layers.
        →xavier_initializer(uniform=False)))
    model.add(layers.BatchNormalization()) # batch norm
    model.add(layers.MaxPooling2D((2, 2))) # max pooling
    model.add(layers.Flatten()) # flatten
    model.add(layers.Dense(784,
        activation='relu',
        kernel_regularizer=tf.contrib.layers.
        →l2_regularizer(scale=scale))) #L
    # fully-connected 784 w/ ReLu
    model.add(layers.Dense(10)) # fully-connected 10
    model.add(layers.Softmax()) # softmax output

    # compile model w/ Adam optimizer + cross entropy loss
    model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
        loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
        metrics=['accuracy'])

    # train
    history_1 = model.fit(trainData, train_labels,
        validation_data = (validData, valid_labels),
        epochs=epochs,
        batch_size=batch_size,
        callbacks=[TestCallback((testData, test_labels))],
        verbose=0, # 0 = silent, 1 = per epoch
```

```

        shuffle=True)


# display statistics
train_acc, train_loss = history.history["acc"], history.history["loss"]
val_acc, val_loss = history.history["val_acc"], history.history["val_loss"]
test_acc, test_loss = history.history["test_acc"][0], history.
↪history["test_loss"][0]

display_statistics(train_loss=train_loss, train_acc=train_acc,
                  valid_loss=val_loss, valid_acc=val_acc,
                  test_loss=test_loss, test_acc=test_acc)

```



L2 Normalization with 0.01

Training loss: 1.5349 Training acc: 96.81%
 Validation loss: 1.5601 Validation acc: 94.2%
 Testing loss: 1.5696 Testing acc: 93.43%




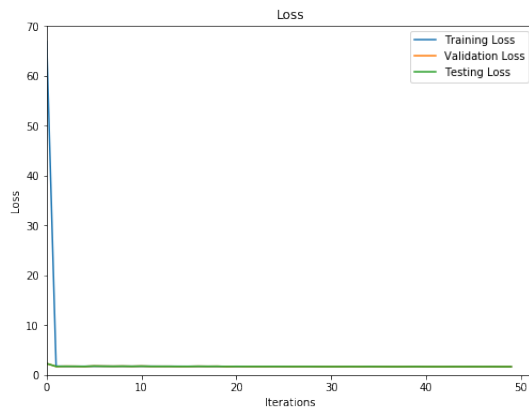
L2 Normalization with 0.1

Training loss: 1.6037 Training acc: 92.77%
 Validation loss: 1.6033 Validation acc: 92.6%
 Testing loss: 1.6058 Testing acc: 92.44%



L2 Normalization with 0.5

Training loss: 1.6729 Training acc: 90.07%
 Validation loss: 1.6705 Validation acc: 90.3%
 Testing loss: 1.6746 Testing acc: 89.83%



2.3.2 Dropout

```
[0]: # training params
learning_rate = 0.0001
epochs = 50
batch_size = 32

# for rate in [0.9, 0.75, 0.5]:
for rate in [0.1, 0.25, 0.5]:
    print("\nDropout with probability {}".format(rate))

    # create model
    model = models.Sequential()
    model.add(layers.InputLayer(input_shape=(28, 28, 1))) # input layer
    model.add(layers.Conv2D(
        filters=32,
        strides=(1, 1),
        kernel_size=[3, 3],
        padding="same",
        activation='relu',
        kernel_initializer=tf.contrib.layers.
    ↪xavier_initializer(uniform=False)))
    model.add(layers.BatchNormalization()) # batch norm
    model.add(layers.MaxPooling2D((2, 2))) # max pooling
    model.add(layers.Flatten()) # flatten
    model.add(layers.Dense(784)) # fully-connected 784
    model.add(layers.Dropout(rate=rate)) # dropout
    model.add(layers.ReLU()) # Relu activation
    model.add(layers.Dense(10)) # fully-connected 10
    model.add(layers.Softmax()) # softmax output

    # compile model w/ Adam optimizer + cross entropy loss
    model.compile(optimizer=tf.keras.optimizers.Adam(lr=learning_rate),
                  loss=tf.keras.losses.CategoricalCrossentropy(from_logits=True),
                  metrics=['accuracy'])

    # train
    history_2 = model.fit(trainData, train_labels,
                          validation_data = (validData, valid_labels),
                          epochs=epochs,
                          batch_size=batch_size,
                          callbacks=[TestCallback((testData, test_labels))],
                          verbose=0, # 0 = silent, 1 = per epoch
                          shuffle=True)

    # display stats
    train_acc, train_loss = history.history["acc"], history.history["loss"]
    val_acc, val_loss = history.history["val_acc"], history.history["val_loss"]
```

```


test_acc, test_loss = history.history["test_acc"][0], history.
↪history["test_loss"][0]

display_statistics(train_loss=train_loss, train_acc=train_acc,
                  valid_loss=val_loss, valid_acc=val_acc,
                  test_loss=test_loss, test_acc=test_acc)

```


Dropout with probability 0.1

Training loss: 1.4831 Training acc: 97.81%
 Validation loss: 1.5229 Validation acc: 93.7%
 Testing loss: 1.5219 Testing acc: 93.98%



Dropout with probability 0.25

Training loss: 1.4794 Training acc: 98.19%
 Validation loss: 1.5156 Validation acc: 94.4%
 Testing loss: 1.527 Testing acc: 93.36%



Dropout with probability 0.5

Training loss: 1.4803 Training acc: 98.09%

Validation loss: 1.5206 Validation acc: 94.0%

Testing loss: 1.5238 Testing acc: 93.69%

